

REDACTPROOF - AI TRANSPARENCY STATEMENT

Generated (UTC): 2026-08-12T00:32:13.033612+00:00 · Platform-level document (applies to every organization)

Model inventory

Detection runs two passes. The deterministic pass (Microsoft Presidio + org-configured regex/dictionary/entity rules) never calls a language model - it's pattern and NER matching, fully explainable per match. The contextual pass calls a large language model only on pages/chunks needing judgment (narrative text, ambiguous entities) and only proposes redactions with a citation and plain-language justification - a human decides on every one before it's ever applied. Production configuration targets Amazon Bedrock (model id set per-deployment via BEDROCK_MODEL_ID, default region us-east-1), under Bedrock's zero-retention configuration - customer content is never used to train Bedrock's underlying models or ours. Every AI-proposed redaction records which model id and prompt version produced it, so any past decision is explainable after the fact, not just at review time.

Data flow



Original content is stored once (per-org encryption key) and never leaves this pipeline: extraction and detection read it, redaction candidates reference spans and boxes (not copies of the source file), export burns approved redactions into a new file and scrubs metadata, and every step writes an immutable, hash-chained audit record. Nothing downstream of the LLM call is allowed to auto-apply a redaction - the human-review step below is a hard gate, not a configurable one.

Human-in-the-loop statement

No document can reach an exported state without at least one human review action. Every AI-proposed redaction is stored as a candidate in a 'suggested' state, never a final decision - a human must approve, reject, or modify each one, with a mandatory exemption code, before export is possible. This is enforced at the API layer (not just the review UI), and every decision is attributed to a specific user and timestamped in the audit trail.

Accuracy report

NOT YET AVAILABLE. The golden-fixture measurement suite this report depends on (specs/05-redaction-pipeline.md: ~30 fixture documents, recall $\geq 95\%$ / precision $\geq 80\%$ targets) has not been built or run yet in this codebase as of this document's generation date. We are not publishing invented numbers here. This section will report actual measured precision/recall against that suite, updated per release, once it exists.

Bias testing summary

NOT YET AVAILABLE, for the same reason as the accuracy report above - no name-ethnicity or demographic-representation testing on detection recall has been conducted yet. This section will report the actual methodology and results once that testing exists, not before.